

MariaDB

The screenshot shows a Microsoft Word document window. The ribbon bar at the top has tabs for Home, Insert, Page Layout, References, Mailings, Review, and View. Below the ribbon are standard toolbar icons for zoom, orientation, page view, read aloud, add text, draw, and highlight. The main content area contains the following text:

Problem Statement:

You work for XYZ Corporation. Their application requires a SQL service that can store data which can be retrieved if required. Implement a suitable RDS engine for the same.

While migrating, you are asked to perform the following tasks:

1. Create a MariaDB Engine based RDS Database.
2. Connect to the DB using the following ways:
 - a. SQL Client for Windows
 - b. Linux based EC2 Instance

Search for RDS service – create database – dashboard.

Amazon RDS

Dashboard

Databases

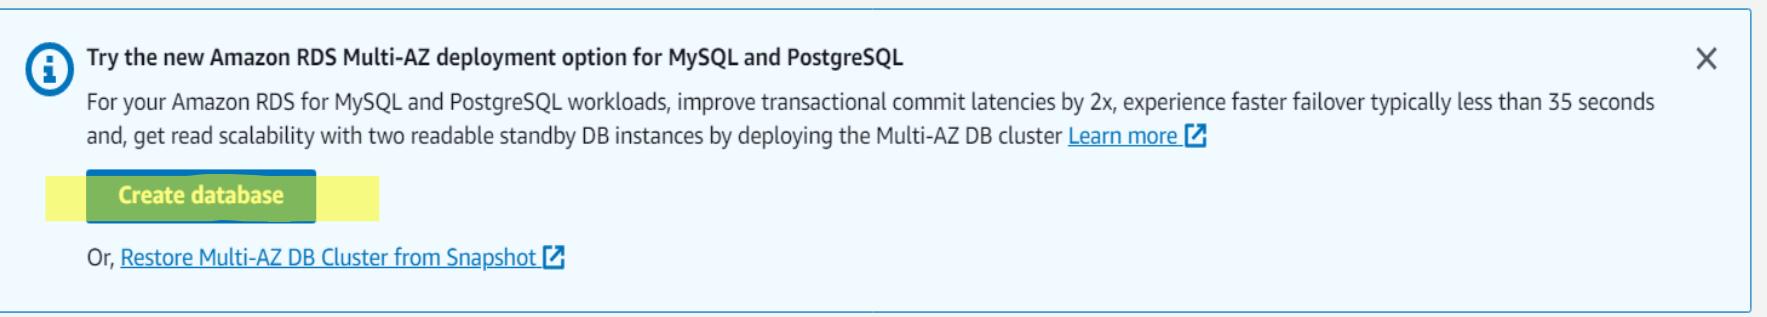
Performance insights

Snapshots

Try the new Amazon RDS Multi-AZ deployment option for MySQL and PostgreSQL
For your Amazon RDS for MySQL and PostgreSQL workloads, improve transactional commit latencies by 2x, experience faster failover typically less than 35 seconds and, get read scalability with two readable standby DB instances by deploying the Multi-AZ DB cluster [Learn more](#)

Create database

Or, [Restore Multi-AZ DB Cluster from Snapshot](#)



Deleting DB instance sid-rds-aws

RDS > Databases

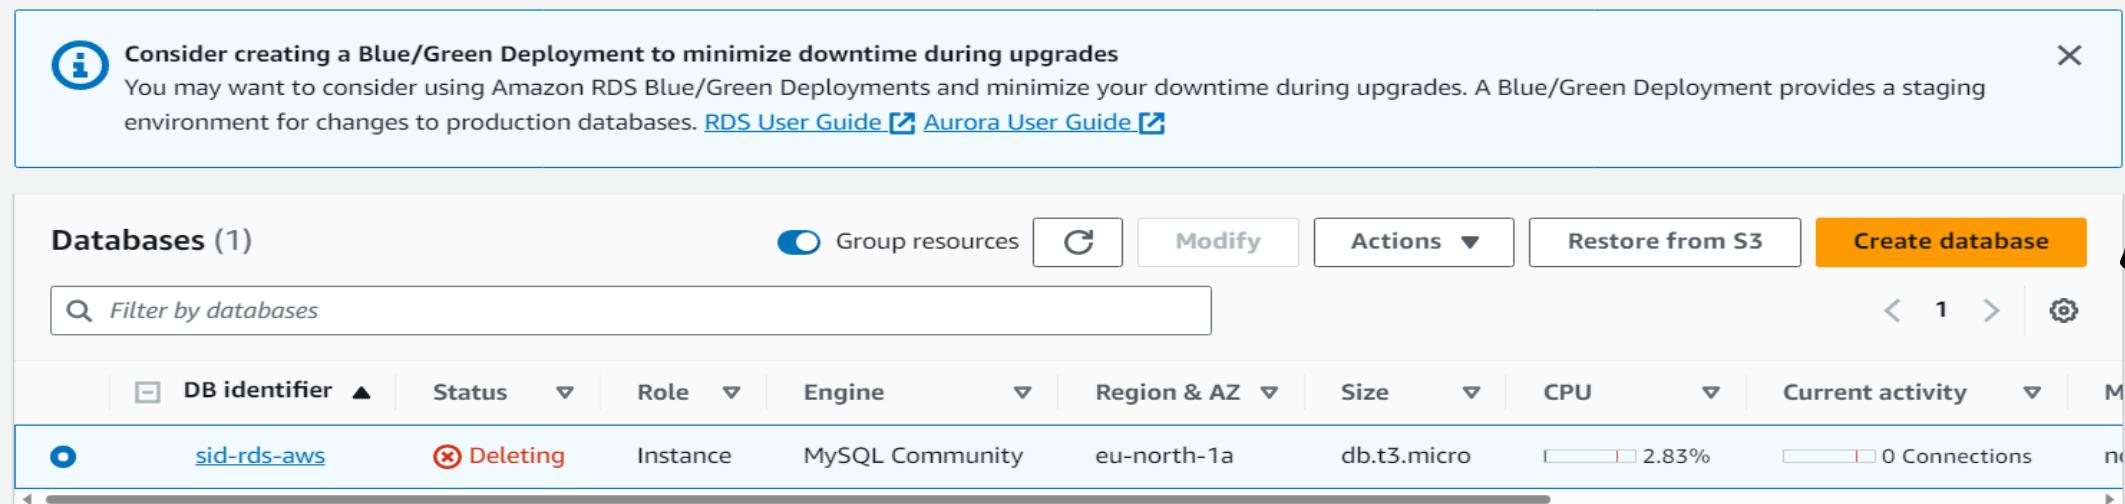
Consider creating a Blue/Green Deployment to minimize downtime during upgrades
You may want to consider using Amazon RDS Blue/Green Deployments and minimize your downtime during upgrades. A Blue/Green Deployment provides a staging environment for changes to production databases. [RDS User Guide](#) [Aurora User Guide](#)

Databases (1)

Group resources [C](#) [Modify](#) [Actions ▾](#) [Restore from S3](#) **Create database**

Filter by databases

DB identifier	Status	Role	Engine	Region & AZ	Size	CPU	Current activity	M
sid-rds-aws	✖ Deleting	Instance	MySQL Community	eu-north-1a	db.t3.micro	2.83%	0 Connections	no



Create database

Choose a database creation method [Info](#)

Standard create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

Aurora (MySQL Compatible)



Aurora (PostgreSQL Compatible)



MySQL



MariaDB



PostgreSQL



Oracle



Microsoft SQL Server



Hide filters

Show versions that support the Amazon RDS Optimized Writes [Info](#)

Amazon RDS Optimized Writes improves write throughput by up to 2x at no additional cost.

Engine Version

MariaDB 10.6.14

Templates

Choose a sample template to meet your use case.

Production

Use defaults for high availability and fast, consistent performance.

Dev/Test

This instance is intended for development use outside of a production environment.

Free tier

Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS.

Settings

DB instance identifier [Info](#)

Type a name for your DB instance. The name must be unique across all DB instances owned by your AWS account in the current AWS Region.

MariaDB-RDS

The DB instance identifier is case-insensitive, but is stored as all lowercase (as in "mydbinstance"). Constraints: 1 to 60 alphanumeric characters or hyphens. First character must be a letter. Can't contain two consecutive hyphens. Can't end with a hyphen.

Credentials Settings

Master username [Info](#)

Type a login ID for the master user of your DB instance.

admin

1 to 16 alphanumeric characters. The first character must be a letter.

Manage master credentials in AWS Secrets Manager

Manage master user credentials in Secrets Manager. RDS can generate a password for you and manage it throughout its lifecycle.

If you manage the master user credentials in Secrets Manager, some RDS features aren't supported.

[Learn more](#)

Auto generate a password

Amazon RDS can generate a password for you, or you can specify your own password.

Master password [Info](#)

Constraints: At least 8 printable ASCII characters. Can't contain any of the following: / (slash), ' (single quote), " (double quote) and @ (at sign).

Confirm master password [Info](#)

Instance configuration

The DB instance configuration options below are limited to those supported by the engine that you selected above.

DB instance class [Info](#)

▼ Hide filters

Show instance classes that support Amazon RDS Optimized Writes [Info](#)

Amazon RDS Optimized Writes improves write throughput by up to 2x at no additional cost.

Include previous generation classes

Standard classes (includes m classes)

Memory optimized classes (includes r and x classes)

Burstable classes (includes t classes)

db.t3.micro

2 vCPUs 1 GiB RAM Network: 2,085 Mbps



Storage

Storage type [Info](#)

General Purpose SSD (gp2)

Baseline performance determined by volume size



Allocated storage [Info](#)

20



The minimum value is 20 GiB and the maximum value is 6,144 GiB

i After you modify the storage for a DB instance, the status of the DB instance will be in storage-optimization. Your instance will remain available as the storage-optimization operation completes.

[Learn more](#)

Storage autoscaling [Info](#)

Provides dynamic scaling support for your database's storage based on your application's needs.

Enable storage autoscaling

Enabling this feature will allow the storage to increase after the specified threshold is exceeded.

Availability & durability

Multi-AZ deployment [Info](#)

Create a standby instance (recommended for production usage)

Creates a standby in a different Availability Zone (AZ) to provide data redundancy, eliminate I/O freezes, and minimize latency spikes during system backups.

Do not create a standby instance

Connectivity [Info](#)

Public access [Info](#)

Yes

RDS assigns a public IP address to the database. Amazon EC2 instances and other resources outside of the VPC can connect to your database. Resources inside the VPC can also connect to the database. Choose one or more VPC security groups that specify which resources can connect to the database.

No

RDS doesn't assign a public IP address to the database. Only Amazon EC2 instances and other resources inside the VPC can connect to your database. Choose one or more VPC security groups that specify which resources can connect to the database.

VPC security group (firewall) [Info](#)

Choose one or more VPC security groups to allow access to your database. Make sure that the security group rules allow the appropriate incoming traffic.

Choose existing

Choose existing VPC security groups

Create new

Create new VPC security group

Existing VPC security groups

Choose one or more options

default X



Availability Zone [Info](#)

No preference



RDS Proxy

RDS Proxy is a fully managed, highly available database proxy that improves application scalability, resiliency, and security.

Create an RDS Proxy [Info](#)

RDS automatically creates an IAM role and a Secrets Manager secret for the proxy. RDS Proxy has additional costs. For more information, see [Amazon RDS Proxy pricing](#).

Certificate authority - optional [Info](#)

Using a server certificate provides an extra layer of security by validating that the connection is being made to an Amazon database. It does so by checking the server certificate that is automatically installed on all databases that you provision.

rds-ca-2019 (default)

Expiry: Aug 22, 2024



If you don't select a certificate authority, RDS chooses one for you.

► Additional configuration

Database authentication

Database authentication options [Info](#)

Password authentication

Authenticates using database passwords.

Password and IAM database authentication

Authenticates using the database password and user credentials through AWS IAM users and roles.

DB subnet group [Info](#)

Choose the DB subnet group. The DB subnet group defines which subnets and IP ranges the DB instance can use in the VPC that you selected.

default-vpc-080f6a5ebfa0d9944

3 Subnets, 3 Availability Zones



Monitoring

Enable Enhanced monitoring

Enabling Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU.

▼ Additional configuration

Database options, encryption turned off, backup turned off, backtrack turned off, maintenance, CloudWatch Logs, delete protection turned off.

Database options

Initial database name [Info](#)

MariaDB

If you do not specify a database name, Amazon RDS does not create a database.

DB parameter group [Info](#)

default.mariadb10.6

Option group [Info](#)

default:mariadb-10-6

Backup

Enable automated backups

Creates a point-in-time snapshot of your database

Encryption

Enable encryption

Choose to encrypt the given instance. Master key IDs and aliases appear in the list after they have been created using the AWS Key Management Service console. [Info](#)

Log exports

Select the log types to publish to Amazon CloudWatch Logs

Audit log

Error log

General log

Slow query log

Maintenance window [Info](#)

Select the period you want pending modifications or maintenance applied to the database by Amazon RDS.

Choose a window

No preference

Deletion protection

Enable deletion protection

Protects the database from being deleted accidentally. While this option is enabled, you can't delete the database.

Estimated Monthly costs

DB instance

12.41 USD

Storage

2.40 USD

Total

14.81 USD

This billing estimate is based on on-demand usage as described in [Amazon RDS Pricing](#). Estimate does not include costs for backup storage, IOs (if applicable), or data transfer.

Estimate your monthly costs for the DB Instance using the [AWS Simple Monthly Calculator](#).

Estimated monthly costs

The Amazon RDS Free Tier is available to you for 12 months. Each calendar month, the free tier will allow you to use the Amazon RDS resources listed below for free:

- 750 hrs of Amazon RDS in a Single-AZ db.t2.micro, db.t3.micro or db.t4g.micro Instance.
- 20 GB of General Purpose Storage (SSD).
- 20 GB for automated backup storage and any user-initiated DB Snapshots.

[Learn more about AWS Free Tier](#).

When your free usage expires or if your application use exceeds the free usage tiers, you simply pay standard, pay-as-you-go service rates as described in the [Amazon RDS Pricing page](#).

 You are responsible for ensuring that you have all of the necessary rights for any third-party products or services that you use with AWS services.

Cancel

Create database

Successfully deleted DB instance sid-rds-aws

Successfully created database mariadb-rds

You can use settings from mariadb-rds to simplify configuration of suggested database add-ons while we finish creating your DB for you.

[View connection details](#)

RDS > Databases



Consider creating a Blue/Green Deployment to minimize downtime during upgrades

You may want to consider using Amazon RDS Blue/Green Deployments and minimize your downtime during upgrades. A Blue/Green Deployment provides a staging environment for changes to production databases. [RDS User Guide](#) [Aurora User Guide](#)

Databases (1)

Group resources



Modify

Actions ▾

Restore from S3

Create database

< 1 >

DB identifier	Status	Role	Engine	Region & AZ	Size	CPU	Current activity	Maintenance
mariadb-rds	Available	Instance	MariaDB	eu-north-1b	db.t3.micro	-	none	

Connectivity & security

Monitoring

Logs & events

Configuration

Maintenance & backups

Tags

Connectivity & security

Endpoint & port

Endpoint
mariadb-rds.cf0cy3osjtej.eu-north-1.rds.amazonaws.com

Port
3306

Networking

Availability Zone
eu-north-1b

VPC
[vpc-080f6a5ebfa0d9944](#)

Subnet group
[subnets-00000000000000000000](#)

Security

VPC security groups
[launch-wizard-17 \(sg-0ea6c1cd698da37c8\)](#)
 Active
[default \(sg-0f4e658f75a245dcd\)](#)
 Active

Publicly accessible

Copy the endpoint using command – , mysql –h [endpoint] –u _ -p _ ,you can connect to the instance but before that give these command sin instance for providing the path for connect to RDS.

1] sudo add-apt-repository -y ppa:ondrej/php

2] sudo apt install php5.6 mysql-client php5.6-mysqli –y .

```
root@ip-172-31-40-201:/home/ubuntu# sudo mysql -h mariadb-rds.cf0cy3osjtej.eu-north-1.rds.amazonaws.com -u admin -p
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 41
Server version: 5.5.5-10.6.14-MariaDB managed by https://aws.amazon.com/rds/

Copyright (c) 2000, 2023, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> show databases;
```

```
| Database           |
+-----+
| MariaDB            |
| information_schema |
| innodb             |
| mysql              |
| performance_schema |
| sys                |
+-----+
6 rows in set (0.00 sec)

mysql> use MariaDB;
Database changed
mysql> create table data(firstname varchar(15), email varchar(25))
   -> select * from data;
ERROR 1146 (42S02): Table 'MariaDB.data' doesn't exist
mysql> create table data(eid int);
Query OK, 0 rows affected (0.00 sec)

mysql> show tables;
+-----+
| Tables_in_MariaDB |
+-----+
| data              |
+-----+
1 row in set (0.00 sec)

mysql> []
```

i-057f550941b4ed27d (mariadb_instance)
Public IPs: 13.51.197.192 Private IPs: 172.31.40.201

Now it's done for ec2 instance next we are going to the same for windows user for this purpose we already install the xamapp .

Download and install the XAMAPP choose MYSQL action start package for providing the path and then go to shell.



Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
MariaDB [(none)]> show databases;
+-----+
| Database |
+-----+
| MariaDB   |
| information_schema |
| innodb    |
| mysql     |
| performance_schema |
| sys       |
+-----+
6 rows in set (0.289 sec)

MariaDB [(none)]> use MariDB
ERROR 1049 (42000): Unknown database 'MariDB'
MariaDB [(none)]> use MariaDB;
Database changed
MariaDB [MariaDB]> create table data(eid, int);
ERROR 1064 (42000): You have an error in your SQL syntax; check the manual that corresponds to your MariaDB server version for the right syntax to use near ' int)' at line 1
MariaDB [MariaDB]> create table data(eid int);
ERROR 1050 (42S01): Table 'data' already exists
MariaDB [MariaDB]> show table data
-> Bye
```

We already created table data in ec2 instance using the same RDS that is why table data already exists.

Aurora

The screenshot shows a Microsoft Word document window. The ribbon bar at the top has tabs for Home, Insert, Page Layout, References, Mailings, Review, and View. Below the ribbon are standard toolbar icons for zoom, orientation, and orientation. The main content area contains the following text:

Problem Statement:

You work for XYZ Corporation. Their application requires a SQL service that can store data which can be retrieved if required. Implement a suitable RDS engine for the same.

While migrating, you are asked to perform the following tasks:

1. Create an AuroraDB Engine based RDS Database.
2. Create 2 Read Replicas in different availability zones for better infrastructure availability.

Create database

Choose a database creation method [Info](#)

Standard create

You set all of the configuration options, including ones for availability, security, backups, and maintenance.

Easy create

Use recommended best-practice configurations. Some configuration options can be changed after the database is created.

Engine options

Engine type [Info](#)

Aurora (MySQL Compatible)



Aurora (PostgreSQL Compatible)



MySQL



MariaDB



PostgreSQL



Oracle



Microsoft SQL Server



Engine version [Info](#)

View the engine versions that support the following database features.

▼ Hide filters

Show versions that support the global database feature

Allows a single Amazon Aurora database to span multiple AWS Regions.

Show versions that support the parallel query feature

Improves the performance of analytic queries by pushing processing down to the Aurora storage layer.

Show versions that support Serverless v2

Offers instance scaling for even the most demanding workloads.

Available versions (20/20) [Info](#)

Aurora MySQL 3.04.1 (compatible with MySQL 8.0.28)

Templates

Choose a sample template to meet your use case.

Production

Use defaults for high availability and fast, consistent performance.

Dev/Test

This instance is intended for development use outside of a production environment.

Settings

DB cluster identifier [Info](#)

Enter a name for your DB cluster. The name must be unique across all DB clusters owned by your AWS account in the current AWS Region.

db-aurora-rds

The DB cluster identifier is case-insensitive, but is stored as all lowercase (as in "mydbcluster"). Constraints: 1 to 60 alphanumeric characters or hyphens. First character must be a letter. Can't contain two consecutive hyphens. Can't end with a hyphen.

Credentials Settings

Master username [Info](#)

Type a login ID for the master user of your DB instance.

admin

1 to 32 alphanumeric characters. The first character must be a letter.

Manage master credentials in AWS Secrets Manager

Manage master user credentials in Secrets Manager. RDS can generate a password for you and manage it throughout its lifecycle.

ⓘ If you manage the master user credentials in Secrets Manager, some RDS features aren't supported.

[Learn more](#)

Auto generate a password

Amazon RDS can generate a password for you, or you can specify your own password.

Master password [Info](#)

Constraints: At least 8 printable ASCII characters. Can't contain any of the following: / (slash), '(single quote)', '(double quote)' and @ (at sign).

Confirm master password [Info](#)

Cluster storage configuration - new [Info](#)

Choose the storage configuration for the Aurora DB cluster that best fits your application's price predictability and price performance needs.

Configuration options

Database instance, storage, and I/O charges vary depending on the configuration. [Learn more](#)

Aurora Standard

- Cost-effective pricing for many applications with moderate I/O usage (I/O costs <25% of total database costs).
- Pay-per-request I/O charges apply. DB instance and storage prices don't include I/O usage.

Aurora I/O-Optimized

- Predictable pricing for all applications. Improved price performance for I/O-intensive applications (I/O costs >25% of total database costs).
- No additional charges for read/write I/O operations. DB instance and storage prices include I/O usage.

Availability & durability

Multi-AZ deployment [Info](#)

Create an Aurora Replica or Reader node in a different AZ (recommended for scaled availability)

Creates an Aurora Replica for fast failover and high availability.

Don't create an Aurora Replica

Connectivity [Info](#)



Compute resource

Choose whether to set up a connection to a compute resource for this database. Setting up a connection will automatically change connectivity settings so that the compute resource can connect to this database.

Don't connect to an EC2 compute resource

Don't set up a connection to a compute resource for this database. You can manually set up a connection to a compute resource later.

Connect to an EC2 compute resource

Set up a connection to an EC2 compute resource for this database.

Network type [Info](#)

To use dual-stack mode, make sure that you associate an IPv6 CIDR block with a subnet in the VPC you specify.

IPv4

Your resources can communicate only over the IPv4 addressing protocol.

Dual-stack mode

Your resources can communicate over IPv4, IPv6, or both.

Virtual private cloud (VPC) [Info](#)

Choose the VPC. The VPC defines the virtual networking environment for this DB cluster.

Default VPC (vpc-080f6a5ebfa0d9944)

3 Subnets, 3 Availability Zones

Only VPCs with a corresponding DB subnet group are listed.

ⓘ After a database is created, you can't change its VPC.

DB subnet group [Info](#)

Choose the DB subnet group. The DB subnet group defines which subnets and IP ranges the DB cluster can use in the VPC that you selected.

default-vpc-080f6a5ebfa0d9944

3 Subnets, 3 Availability Zones

Public access [Info](#)

Yes

RDS assigns a public IP address to the cluster. Amazon EC2 instances and other resources outside of the VPC can connect to your cluster. Resources inside the VPC can also connect to the cluster. Choose one or more VPC security groups that specify which resources can connect to the cluster.

No

RDS doesn't assign a public IP address to the cluster. Only Amazon EC2 instances and other resources inside the VPC can connect to your cluster. Choose one or more VPC security groups that specify which resources can connect to the cluster.

VPC security group (firewall) [Info](#)

VPC security group (firewall) [Info](#)

Choose one or more VPC security groups to allow access to your database. Make sure that the security group rules allow the appropriate incoming traffic.

Choose existing

Choose existing VPC security groups

Create new

Create new VPC security group

Existing VPC security groups

Choose one or more options ▾

default X

RDS Proxy

RDS Proxy is a fully managed, highly available database proxy that improves application scalability, resiliency, and security.

Create an RDS Proxy [Info](#)

RDS automatically creates an IAM role and a Secrets Manager secret for the proxy. RDS Proxy has additional costs. For more information, see [Amazon RDS Proxy pricing](#).

Certificate authority - optional [Info](#)

Using a server certificate provides an extra layer of security by validating that the connection is being made to an Amazon database. It does so by checking the server certificate that is automatically installed on all databases that you provision.

rds-ca-2019 (default)

Expiry: Aug 22, 2024

If you don't select a certificate authority, RDS chooses one for you.

► Additional configuration

Read replica write forwarding

Turn on local write forwarding [Info](#)

Issues write operations from reader DB instances within the same DB cluster.

Database authentication [Info](#)

Password authentication is always active for your database engine. You can also turn on additional authentication methods for your database below.

IAM database authentication

Authenticates using IAM database authentication.

Kerberos authentication

Authenticates using Kerberos authentication through an AWS Directory Service for Microsoft Active Directory.

Monitoring

Enable Enhanced monitoring

Enabling Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU.

▼ Additional configuration

Database options, encryption turned off, failover, backup turned on, backtrack turned off, maintenance, CloudWatch Logs, delete protection turned off.

Database options

Initial database name [Info](#)

MyAurora

If you do not specify a database name, Amazon RDS does not create a database.

DB cluster parameter group [Info](#)

default.aurora-mysql8.0

DB parameter group [Info](#)

default.aurora-mysql8.0

Option group [Info](#)

default:aurora-mysql-8-0

Failover priority

No preference

Backup

Backup retention period [Info](#)

The number of days (1-35) for which automatic backups are kept.

1 ▾ day

Copy tags to snapshots

Encryption

Enable encryption

Choose to encrypt the given instance. Master key IDs and aliases appear in the list after they have been created using the AWS Key Management Service console. [Info](#)

Log exports

Select the log types to publish to Amazon CloudWatch Logs

- Audit log
- Error log
- General log
- Slow query log

IAM role

The following service-linked role is used for publishing logs to CloudWatch Logs.

RDS service-linked role

Maintenance

Auto minor version upgrade [Info](#)

Enable auto minor version upgrade

Enabling auto minor version upgrade will automatically upgrade to new minor versions as they are released. The automatic upgrades occur during the maintenance window for the database.

Maintenance window [Info](#)

Select the period you want pending modifications or maintenance applied to the database by Amazon RDS.

Choose a window

No preference

Deletion protection

Enable deletion protection

Protects the database from being deleted accidentally. While this option is enabled, you can't delete the database.

If connection not estb go on VPC security group choose all traffic allowed security group.

You can use settings from db-aurora-rds to simplify configuration of [suggested database add-ons](#) while we finish creating your DB for you.

RDS > Databases



Consider creating a Blue/Green Deployment to minimize downtime during upgrades

You may want to consider using Amazon RDS Blue/Green Deployments and minimize your downtime during upgrades. A Blue/Green Deployment provides a staging environment for changes to production databases. [RDS User Guide](#) [Aurora User Guide](#)

Databases (3)

Group resources



Modify

Actions ▾

Restore from S3

Create database

< 1 >



Filter by databases

DB identifier	Status	Role	Engine	Region & AZ	Size	CPU	Current
db-aurora-rds	Available	Regional cluster	Aurora MySQL	eu-north-1	2 instances	-	-
db-aurora-rds-instance-1	Creating	Writer instance	Aurora MySQL	eu-north-1b	db.t3.medium	-	-
db-aurora-rds-instance-1-eu-north-1a	Creating	Reader instance	Aurora MySQL	eu-north-1a	db.t3.medium	-	-

For modifying go in modify do the changes and apply immediately.

db-aurora-rds	Available	Regional cluster	Aurora MySQL	eu-north-1	2 instances	-
db-aurora-rds-instance-1	Available	Writer instance	Aurora MySQL	eu-north-1b	db.t3.medium	<div style="width: 14.10%;">14.10%</div>
db-aurora-rds-instance-1-eu-north-1a	Available	Reader instance	Aurora MySQL	eu-north-1a	db.t3.medium	<div style="width: 12.53%;">12.53%</div>

Connectivity & security

Monitoring

Logs & events

Configuration

Zero-ETL integrations

Maintenance & backups

Tags

Endpoints (2)

Actions ▾

Create custom endpoint

Find resources

< 1 >



Endpoint name	Status	Type	Port
db-aurora-rds.cluster-cf0cy3osjtej.eu-north-1.rds.amazonaws.com	Available	Writer	3306
db-aurora-rds.cluster-ro-cf0cy3osjtej.eu-north-1.rds.amazonaws.com	Available	Reader	3306

```
# mysql -h db-aurora-rds.cluster-cf0cy3osjtej.eu-north-1.rds.amazonaws.com -u admin -p  
Enter password: *****  
Welcome to the MariaDB monitor. Commands end with ; or \g.  
Your MySQL connection id is 89  
Server version: 8.0.28 Source distribution
```

```
Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.
```

```
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.
```

```
MySQL [(none)]> show databases;  
+-----+  
| Database |  
+-----+  
| MyAurora |  
| information_schema |  
| mysql |  
| performance_schema |  
| sys |  
+-----+  
5 rows in set (0.991 sec)
```

```
MySQL [(none)]> use MyAurora;  
Database changed  
MySQL [MyAurora]> create table data(eid int);  
Query OK, 0 rows affected (0.390 sec)
```

```
MySQL [MyAurora]> show tables;  
+-----+  
| Tables_in_MyAurora |  
+-----+  
| data |  
+-----+  
1 row in set (0.414 sec)
```

```
MySQL [MyAurora]> quit  
Bye
```

```
# mysql -h db-aurora-rds-instance-1-eu-north-1a.cf0cy3osjtej.eu-north-1.rds.amazonaws.com -u admin -p  
Enter password: *****  
Welcome to the MariaDB monitor. Commands end with ; or \g.  
Your MySQL connection id is 54  
Server version: 8.0.28 Source distribution  
Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.  
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.  
MySQL [(none)]> show databases;  
+-----+  
| Database |  
+-----+  
| MyAurora |  
| information_schema |  
| mysql |  
| performance_schema |  
| sys |  
+-----+  
5 rows in set (0.222 sec)  
MySQL [(none)]> use MyAurora;  
Database changed  
MySQL [MyAurora]> create table data(eid int);  
ERROR 1050 (42S01): Table 'data' already exists  
MySQL [MyAurora]> create database mydata;  
ERROR 1836 (HY000): Running in read-only mode  
MySQL [MyAurora]>
```

When we login as reader so we don't have permission to create databases.

RDS > Databases > db-aurora-rds

db-aurora-rds

Related

Filter by databases

DB identifier	Status	Role	Engine	Region & AZ
db-aurora-rds	Available	Regional cluster	Aurora MySQL	eu-north-1
db-aurora-rds-instance-1	Available	Writer instance	Aurora MySQL	eu-north-1b
db-aurora-rds-instance-1-eu-north-1a	Available	Reader instance	Aurora MySQL	eu-north-1a



Modify

Actions ▾

Stop temporarily

Delete

Set up EC2 connection

Set up Lambda connection

Add AWS Region

Add reader

Create cross-Region read replica

Create Blue/Green Deployment - new

Create clone

Promote

RDS > Databases > Add reader

Add reader

You are creating a replica DB instance from a source DB instance. This new DB instance will have the source DB instance's DB security groups and DB parameter groups.

Settings

Aurora replica source

Source DB cluster identifier

db-aurora-rds-instance-1

Role: Writer instance Parent: db-aurora-rds

DB instance identifier

DB instance identifier. This is the unique key that identifies a DB instance. This parameter is stored as a lowercase string (for example, mydbinstance).

my-reader

Instance configuration

The DB instance configuration options below are limited to those supported by the engine that you selected above.

DB instance class [Info](#)

▼ Hide filters

Include previous generation classes

Serverless v2

Memory optimized classes (includes r classes)

Burstable classes (includes t classes)

AWS Region

Destination Region

The Region where the replica will be launched.

EU (Stockholm)



Connectivity

Public access

Publicly accessible
RDS assigns a public IP address to the database. Amazon EC2 instances and other resources outside of the VPC can connect to your database. Resources inside the VPC can also connect to the database. Choose one or more VPC security groups that specify which resources can connect to the database.

Not publicly accessible

No IP address is assigned to the DB instance. EC2 instances and devices outside the VPC can't connect.

Availability Zone [Info](#)

The EC2 Availability Zone that the database will be created in.

eu-north-1c



► Additional configuration

Database authentication [Info](#)

Password authentication is always active for your database engine. You can also turn on additional authentication methods for your database below.

IAM database authentication

Authenticates using IAM database authentication.

Kerberos authentication

Authenticates using Kerberos authentication through an AWS Directory Service for Microsoft Active Directory.

▼ Additional configuration
encryption turned off, failover, Enhanced Monitoring turned off, maintenance, CloudWatch Logs

DB parameter group [Info](#)
default.aurora-mysql8.0

Failover priority
No preference

Encryption

Enable encryption
Choose to encrypt the given instance. Master key IDs and aliases appear in the list after they have been created using the AWS Key Management Service console. [Info](#)

Enable Enhanced monitoring
Enabling Enhanced monitoring metrics are useful when you want to see how different processes or threads use the CPU.

Maintenance

Auto minor version upgrade [Info](#)

Enable auto minor version upgrade
Enabling auto minor version upgrade will automatically upgrade to new minor versions as they are released. The automatic upgrades occur during the maintenance window for the database.

[Cancel](#) [Add reader](#)

Databases (4) (Group resources C [Modify](#) [Actions ▾](#) [Restore from S3](#) [Create database](#)

Q Filter by databases < 1 >

DB identifier	Status	Role	Engine	Region & AZ	Size	CPU
● db-aurora-rds	✔ Available	Regional cluster	Aurora MySQL	eu-north-1	3 instances	-
○ db-aurora-rds-instance-1	✔ Available	Writer instance	Aurora MySQL	eu-north-1b	db.t3.medium	<div style="width: 11.74%;">11.74%</div>
○ db-aurora-rds-instance-1-eu-north-1a	✔ Available	Reader instance	Aurora MySQL	eu-north-1a	db.t3.medium	<div style="width: 10.42%;">10.42%</div>
○ my-reader	✔ Available	Reader instance	Aurora MySQL	eu-north-1c	db.t3.medium	-

DynamoDB

Problem Statement:

You work for XYZ Corporation. Their application requires a database service that can store data which can be retrieved if required. Implement a suitable service for the same.

While migrating, you are asked to perform the following tasks:

1. Create a DynamoDB table with partition key as ID.
2. Add 5 items to the DynamoDB table.
3. Take backup and delete the table.

Search for DynamoDB service and go for create tables.

DynamoDB



Share your feedback on Amazon DynamoDB

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how we're doing.

Share feedback



Dashboard

Tables

Update settings

Explore items

PartiQL editor

Backups

Exports to S3

Imports from S3

Reserved capacity

Settings

Database

Amazon DynamoDB

A fast and flexible NoSQL database service for any scale

DynamoDB is a fully managed, key-value, and document database that delivers single-digit-millisecond performance at any scale.

Get started

Create a new table to start exploring DynamoDB.

Create table

DynamoDB > Tables > Create table

Create table

Table details Info

DynamoDB is a schemaless database that requires only a table name and a primary key when you create the table.

Table name

This will be used to identify your table.

DB_table

Between 3 and 255 characters, containing only letters, numbers, underscores (_), hyphens (-), and periods (.)

Partition key

The partition key is part of the table's primary key. It is a hash value that is used to retrieve items from your table and allocate data across hosts for scalability and availability.

Reg

Number

1 to 255 characters and case sensitive.

Sort key - optional

You can use a sort key as the second part of a table's primary key. The sort key allows you to sort or search among all items sharing the same partition key.

Enter the sort key name

String

1 to 255 characters and case sensitive.

Table settings

Default settings

The fastest way to create your table. You can modify these settings now or after your table has been created.

Customize settings

Use these advanced features to make DynamoDB work better for your needs.

Default table settings

These are the default settings for your new table. You can change some of these settings after creating the table.

Setting	Value	Editable after creation
Table class	DynamoDB Standard	Yes
Capacity mode	Provisioned	Yes
Provisioned read capacity	5 RCU	Yes
Provisioned write capacity	5 WCU	Yes
Auto scaling	On	Yes
Local secondary indexes	-	No
Global secondary indexes	-	Yes
Encryption key management	Owned by Amazon DynamoDB	Yes
Deletion protection	Off	Yes

Tags

Tags are pairs of keys and optional values, that you can assign to AWS resources. You can use tags to control access to your resources or track your AWS spending.

No tags are associated with the resource.

Add new tag

You can add 50 more tags.

Cancel

Create table

✓ The DB_table table was created successfully.

X

DynamoDB > Tables

Tables (1) Info									
<input type="checkbox"/>	Name	Status	Partition key	Sort key	Indexes	Deletion protection	Read capacity mode	Write capacity mode	Total
<input type="checkbox"/>	DB_table	Active	Reg (N)	-	0	Off	Provisioned (5)	Provisioned (5)	0 bytes

DynamoDB

X

ⓘ Share your feedback on Amazon DynamoDB

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how we're doing.

Share feedback

ⓘ

Dashboard

Tables

Update settings

Explore items

PartiQL editor

Backups

Exports to S3

Imports from S3

Reserved capacity

Settings

Tables (1)	
<input type="checkbox"/>	Any tag key
<input type="checkbox"/>	Any tag value
<input type="checkbox"/>	<input type="text"/> Find tables by table name
<input type="checkbox"/>	DB_table

DB_table

Autopreview

[View table details](#)

▼ Scan or query items

Scan

Query

Select a table or index

Table - DB_table

Select attribute projection

All attributes

► Filters

[Run](#)

Reset

✓ Completed. Read capacity units consumed: 0.5

Items returned (0)

Actions ▾

Create item

< 1 >

X

No items

No items to display.

[Create item](#)

You can add items by going into explore items and click table and choose explore table item.

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how we're doing.

DynamoDB > Explore items: DB_table > Create item

Create item

Form

JSON view

You can add, remove, or edit the attributes of an item. You can nest attributes inside other attributes up to 32 levels deep. [Learn more](#)

Attributes		
Attribute name	Value	Type
Reg - Partition key	1	Number
Name	Sid	String

[Add new attribute ▾](#)

[Cancel](#) [Create item](#)

Go on add new attribute- enter the value- create item.

Items returned (5)

[C](#) [Actions ▾](#) [Create item](#)

< 1 > | [⚙️](#) [✖️](#)

	Reg (Number)	Address	Class	Class	Name	School	State
4		Newcity	7	7		AIPS	California
5		Newcity	5	5	tobi	AIPS	
3			4	4		AIPS	
2			3	3	itachi		
1					Sid		

To create the go on backup option directly.

DynamoDB

X

Dashboard

Tables

Update settings

Explore items

PartiQL editor

Backups

Exports to S3

Imports from S3

Reserved capacity

Settings

Share your feedback on Amazon DynamoDB
Your feedback is an important part of helping us provide a better customer experience. Take this short survey.

DynamoDB > Tables

Tables (1/1) [Info](#)

Find tables by table name

Any tag key

<input checked="" type="checkbox"/>	Name ▲	Status	Partition key	Sort key	Indexes	Deletion prot
<input checked="" type="checkbox"/>	DB_table	<input checked="" type="checkbox"/> Active	Reg (N)	-	0	<input checked="" type="checkbox"/> Off

Backups [Info](#)

Backup settings [Info](#)

Settings apply to new backups in this account and Region.

[Turn off](#)

Advanced features with AWS Backup

Activated

Allow options for cross-Region and cross-account copy, cost allocation tags, and cold storage tiering for backups.

Backups (0) [Info](#)

[Schedule automatic backups](#) and [view backup job details](#) in [AWS Backup](#)



[View details](#)

[Restore](#)

[Copy](#)

[Delete](#)

[Create backup ▲](#)

[Create on-demand backup](#)

[Schedule backups](#)

Find backups by ARN or name

Name

▼

Table

▼

Status

▼

Creation time

▼

ARN

▼

Size

▼

No backups

Create a backup to save your data.

[Create backup](#)

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how.

Share feedback

Create on-demand backup

Create a one-time snapshot backup of your table. Schedule automatic backups of your table in [AWS Backup](#)

Source table [Info](#)

Source table

DB_table



Backup settings [Info](#)

Default settings

Create a backup that stays in warm storage.

Customize settings

Create a backup that can transition to cold storage and be deleted as it ages.

Backup management [Info](#)

Backup with AWS Backup

Creates a backup with AWS Backup encryption and ARN. Includes options for cross-Region and cross-account copy, tags and cold storage.

Backup with DynamoDB

Creates a backup with DynamoDB encryption and ARN. Additional features not supported.

Backup name

This will be used to identify your backup.

mydb_backup

Between 3 and 255 characters in length. Only A-Z, a-z, 0-9, underscore characters, hyphens, and periods are allowed.

Cancel

Create backup

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how we're doing.

The mydb_backup backup has been created successfully.

Backups [Info](#)

Backup settings [Info](#)

Settings apply to new backups in this account and Region.

Turn off

Advanced features with AWS Backup

Activated

Allow options for cross-Region and cross-account copy, cost allocation tags, and cold storage tiering for backups.

Backups (1) [Info](#)



View details

Restore

Copy

Delete

Create backup ▾

[Schedule automatic backups](#) and [view backup job details](#) in [AWS Backup](#)

Find backups by ARN or name

< 1 >

	Name	Table	Status	Creation time	ARN	Size
<input type="checkbox"/>	mydb_backup	DB_table	Available	November 27, ...	arn:aws:dynamodb:eu-n	0 bytes

Delete table

Delete table **DB_table** in Europe (Stockholm) permanently? This action cannot be undone.

⚠ Proceeding with this action will delete the table and you won't be able to retrieve this data.

Delete all CloudWatch alarms for **DB_table**.

Create an on-demand backup of **DB_table** before deletion.

You can create an on-demand backup of your table for long-term retention and data archiving. You can then use this backup to restore your data to its exact state before table deletion. Additional charges apply for on-demand backup and restore. For more information see [Pricing](#).

To avoid unintentional deletions, we ask you to provide additional confirmation.

To confirm this deletion, type "confirm".

confirm

Cancel

Delete

Backups (2) [Info](#)



[View details](#)

[Restore](#)

[Copy](#)

[Delete](#)

[Create backup ▾](#)

[Schedule automatic backups](#) and [view backup job details](#) in [AWS Backup](#)

Find backups by ARN or name

< 1 >

<input type="checkbox"/>	Name	Table	Status	Creation time	ARN	Size
<input type="checkbox"/>	DB_table_backup	✓ DB_table	Available	November 27, 2...	arn:aws:dynamodb:eu-n	0 bytes
<input type="checkbox"/>	mydb_backup	✓ DB_table	Available	November 27, 2...	arn:aws:dynamodb:eu-n	0 bytes

We can also create backup while deleting the table.

We can also restore the table using backup.

Backups (1/2) [Info](#)

Schedule automatic backups [\[?\]](#) and view backup job details [\[?\]](#) in AWS Backup [\[?\]](#)

View details

Find backups by ARN or name

Name	Table	Status	Creation time	ARN	Size
DB_table_backup	DB_table	Available	November 27, 2...	arn:aws:dynamodb:eu-n...	0 bytes
<input checked="" type="checkbox"/> mydb_backup	DB_table	Available	November 27, 2...	arn:aws:dynamodb:eu-n...	0 bytes

Completed. Read capacity units consumed: 3.5

Items returned (5)

Actions [\[?\]](#) Create item

Reg (Number)	Address	Class	Name	School	State
3		4		AIPS	
2		3	itachi		
4	Newcity	7		AIPS	Calif...
1			Sid		
5	Newcity	5	tobi	AIPS	

My table and data both restored.

Share your feedback on Amazon DynamoDB

Your feedback is an important part of helping us provide a better customer experience. Take this short survey to let us know how

DynamoDB > Backups > Restore table from backup

Restore table from backup [Info](#)

Restoring a table from a backup will restore it as a new table.

Restore settings

Name of restored table

This name will identify your restored table.

mydb_table2

Between 3 and 255 characters in length. Only A-Z, a-z, 0-9, underscore characters, hyphens, and periods allowed.

Secondary indexes

Restore the entire table

Your restored table will include all local and global secondary indexes.

Restore the table without secondary indexes

Your restored table will exclude all local and global secondary indexes. Restoring this way can be faster and more cost efficient.

Destination AWS Region

Same Region (Stockholm)

Restore the table to the same Region as the original table.

Cross-Region

Restore the table to a different Region for greater redundancy but with higher data transfer costs.

Encryption at rest settings [Info](#)

All user data stored in Amazon DynamoDB is fully encrypted at rest. By default, Amazon DynamoDB manages the encryption key, and you are not charged any fee for using it.

Encryption key management

Owned by Amazon DynamoDB

The AWS KMS key is owned and managed by DynamoDB. You are not charged an additional fee for using this key.

AWS managed key

Key alias: aws/dynamodb. The key is stored in your account and is managed by AWS Key Management Service (AWS KMS). AWS KMS charges apply.

Stored in your account, and owned and managed by you

Choose a key that is owned and managed by you, and stored in AWS KMS.

Restores can take several hours to complete. After your table is restored from the backup, you might need to reapply configuration settings. [Learn more \[?\]](#)

Cancel

Restore

Redshift

Problem Statement:

You work for XYZ Corporation. Their application requires a database service that can store data which can be retrieved if required. Implement suitable service for the same.

While migrating, you are asked to perform the following tasks:

1. Create a Redshift data warehouse.
2. Using the query editor:
 - a. Load some data
 - b. Query the data

Buckets (1) [Info](#)

Buckets are containers for data stored in S3. [Learn more](#) 



[Copy ARN](#)

[Empty](#)

[Delete](#)

[Create bucket](#)

 Find buckets by name

< 1 > 

Name	AWS Region	Access	Creation date
redshi-bucket	Europe (Stockholm) eu-north-1	Objects can be public	November 28, 2023, 15:03:53 (UTC+05:30)

For adding the data into Redshift we first have create IAM role and S3 bucket with csv file.

[IAM](#) > [Roles](#) > Create role

Step 1 Select trusted entity

Step 2 Add permissions

Step 3 Name, review, and create

✓ Select trusted entity

Trusted entity type

- AWS service Allow AWS services like EC2, Lambda, or others to perform actions in this account.
- AWS account Allow entities in other AWS accounts belonging to you or a 3rd party to perform actions in this account.
- Web identity Allows users federated by the specified external web identity provider to assume this role to perform actions in this account.
- SAML 2.0 federation Allow users federated with SAML 2.0 from a corporate directory to perform actions in this account.
- Custom trust policy Create a custom trust policy to enable others to perform actions in this account.

✓ Use case

Allow an AWS service like EC2, Lambda, or others to perform actions in this account.

Service or use case

Choose a use case for the specified service.

Use case

- Redshift - Customizable Allows Redshift clusters to call AWS services on your behalf.
- Redshift Allows Redshift clusters to call AWS services on your behalf.
- Redshift - Scheduler Allow Redshift Scheduler to call Redshift on your behalf.

✓ Next 

S3

All types



9 matches



1



Policy name

▲ Type

▼ Description

<input type="checkbox"/>	AmazonDMSRedshiftS3Role	AWS managed	Provides access to manage S3 settings...
--------------------------	-------------------------	-------------	--

<input checked="" type="checkbox"/>	AmazonS3FullAccess	AWS managed	Provides full access to all buckets via t...
-------------------------------------	--------------------	-------------	--

Role details

Role name

Enter a meaningful name to identify this role.

 Redshift_role_aws ✓

Maximum 64 characters. Use alphanumeric and '+,-,.,@-_-' characters.

Description

Add a short explanation for this role.

 Allows Redshift clusters to call AWS services on your behalf.

Maximum 1000 characters. Use alphanumeric and '+,-,.,@-_-' characters.

Step 1: Select trusted entities

[Edit](#)

Trust policy

```
1  [{}]
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Action": [
7         "sts:AssumeRole"
8       ],
9       "Principal": {
10        "Service": [
11          "redshift.amazonaws.com"
12        ]
13      }
14    }
15  ]
16 ]
```

Step 2: Add permissions

[Edit](#)

Permissions policy summary

Policy name

▲ Type

▼ Attached as

 AmazonS3FullAccess

AWS managed

Permissions policy

Step 3: Add tags

Add tags - optional [Info](#)

Tags are key-value pairs that you can add to AWS resources to help identify, organize, or search for resources.

No tags associated with the resource.

[Add new tag](#)

You can add up to 50 more tags.

[Cancel](#)[Previous](#)[Create role](#)

Role Redshift_role_aws created.

View role **X** **i**

Roles (16) [Info](#)

An IAM role is an identity you can create that has specific permissions with credentials that are valid for short durations. Roles can be assumed by entities that you trust.

Search **C** **Delete** **Create role**

<input type="checkbox"/>	Role name	Trusted entities	Last activity
<input type="checkbox"/>	AWSServiceRoleForRDS	AWS Service: rds (Service-Linked Role)	52 minutes ago
<input type="checkbox"/>	AWSServiceRoleForSupport	AWS Service: support (Service-Linked Role)	25 days ago
<input type="checkbox"/>	AWSServiceRoleForTrustedAdvisor	AWS Service: trustedadvisor (Service-Linked Role)	-
<input type="checkbox"/>	rds-monitoring-role	AWS Service: monitoring.rds	-
<input type="checkbox"/>	Redshift_role_aws ✓	AWS Service: redshift	-
<input type="checkbox"/>	s3crr_role_for_siddhartha-s3crr-source	AWS Service: s3	39 days ago

Search for redshift service – go for Redshift serverless.

Amazon Redshift **X**

Redshift Serverless [New](#)

Provisioned clusters dashboard

Clusters

- Reserved nodes
- Snapshots

Query editor

Query editor v2 [New](#)

Queries and loads

Datasources

Zero-ETL integrations [New](#)

▼ Connect to Redshift clusters

Query data using Redshift query editor

Use the query editor v2 to run queries in your Redshift cluster.

Query data

Work with your client tools

You can connect to Amazon Redshift from your client tools, such as SQL clients, business intelligence (BI) tools, and extract, transform, load (ETL) tools, using JDBC or ODBC drivers.

Cluster

Cluster identifier

Choose your JDBC or ODBC driver

Use JDBC or ODBC drivers to connect to Amazon Redshift from your client tools, such as SQL clients, BI tools, and ETL tools. We recommend using the new Amazon Redshift-specific drivers for better performance and scalability.

Driver

JDBC 4.2 without AWS SDK (.jar)

Clusters (0) [Info](#)

C **Query data** **Actions** **Create cluster**

Filter clusters by property or value

<input type="checkbox"/>	Cluster	Status	Cluster namespace	Availability Zone	Multi-AZ	Storage capacity us...
--------------------------	---------	--------	-------------------	-------------------	----------	------------------------

NEW: Amazon Redshift now supports zero-ETL integration with Amazon Aurora MySQL. Learn how you can get started applying near-real time analytics and machine learning on your transactional data today. [Learn more about Zero-ETL integrations](#)

Try new Amazon Redshift features in preview
Create a workgroup with preview features. Production use of the workgroup is not supported. Use this workgroup for testing only.

[Create preview workgroup](#)

Introducing price-performance targets
Try price-performance targets to improve performance while optimizing costs.

[Learn more](#)

Amazon Redshift Serverless

Serverless dashboard [Info](#)

[G](#) [Query data](#) [Create workgroup](#)

Namespace overview [Info](#)

Namespace data from your account

Total snapshots



Dashboards in my account



Dashboards requiring authorization



Dashboards from other accounts



Dashboards requiring association



Filter namespace

All namespaces [▼](#)



Create workgroup

Workgroup

Workgroup is a collection of compute resources from which an endpoint is created. Compute properties include network and security settings.

Workgroup name

This is a unique name that defines the workgroup.

myfirstredshiftgroup



The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

Capacity [Info](#)

Set the base capacity used to process your data warehouse workloads. The capacity is measured in Redshift processing units (RPUs). To improve query performance, increase the RPU value.

Base capacity

Base RPU capacity is set to 128 RPUs by default. To change the base RPU capacity, choose another value from the list.

128



Range must be 32-512 in increments of 8.

Network and security

Virtual private cloud (VPC)

This VPC defines the virtual networking environment for this database.

vpc-080f6a5ebfa0d9944



VPC security groups

This VPC security group defines which subnets and IP ranges can be used in the VPC.

Choose one or more security groups

sg-0f4e658f75a245dc



The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

Subnet

The subnet in the chosen VPC that is associated with the specified database.

Choose three or more subnet IDs

subnet-06dfdd82fb120796a

subnet-06d7f09cca266494f



subnet-04fef07ed760f4ba5

Enhanced VPC routing

Turning on this option routes network traffic between your serverless database and data repositories through a VPC instead of the internet.

Turn on enhanced VPC routing

Cancel

Next

Choose namespace

Namespace

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

- Create a new namespace
- Add to an existing namespace

Namespace

This is a unique name that defines the namespace.

The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

Database name and password

Database name

The name of the first database in the Amazon Redshift Serverless environment.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

Admin user credentials

IAM credentials provided as your default admin user credentials. To add a new admin username and password, customize admin user credentials.

- Customize admin user credentials

To use the default IAM credentials, clear this option.

Permissions

- ⓘ Associate an IAM role so that your serverless endpoint can LOAD and UNLOAD data. You can create an IAM role as the default for this configuration that has the [AmazonRedshiftAllCommandsFullAccess](#) policy attached. This policy includes permissions to run SQL commands to COPY, UNLOAD, and query data with Amazon Redshift Serverless. This policy also grants permissions to run SELECT statements for related services, such as Amazon S3, Amazon CloudWatch logs, Amazon SageMaker, and AWS Glue. You won't be able to run these SQL commands without an IAM role attached to your namespace.

Now associate IAM role with such permission allowed.

Associate IAM roles

IAM roles

Choose from existing IAM roles. You can associate up to 50 IAM roles.

< 1 >

IAM roles

- [AWSServiceRoleForRedshift](#)
- [aws_red_role](#)
- [Redshift_role_aws](#)

[Cancel](#)

[Associate IAM roles](#)

Permissions policies (3)

You can attach up to 10 managed policies.



[Simulate](#)

Filter by Type

All types

<input type="checkbox"/> Policy name 	Type	Att.
<input type="checkbox"/>  AmazonS3ReadOnlyAccess	AWS managed	1
<input type="checkbox"/>  CloudWatchEventsFullAccess	AWS managed	1
<input type="checkbox"/>  CloudWatchLogsFullAccess	AWS managed	1

Associated IAM roles (1)

Set default ▾

Manage IAM roles ▾

Create, associate, or remove an IAM role. You can associate up to 50 IAM roles. You can also choose an IAM role and set it as the default.

Search for associated IAM role by name, status, or role type

< 1 >

IAM roles ▾

▼ Status

Role type ▾

aws_red_role

Not applied

--

Encryption and security

⚠ Your data is encrypted by default with an AWS owned key. To choose a different key, customize your encryption settings.

Customize encryption settings (advanced)

Audit logging Info

Collects logging information for the database.

Export these logs:

User log

Connection log

User activity log

Cancel

Previous

Next

Review and create

Step 1: Create workgroup

Edit

Workgroup

Workgroup is a collection of compute resources from which an endpoint is created. Compute properties include network and security settings.

Workgroup name
myfirstredshiftgroup

Capacity

Set the base capacity used to process your data warehouse workloads. The capacity is measured in Redshift processing units (RPUs). To improve query performance, increase the RPU value.

Base capacity
128

Network and security

Virtual private cloud (VPC)
[vpc-080f6a5ebfa0d9944](#)

VPC security group
[sg-0f4e658f75a245dcd](#)

Subnet
subnet-06dfdd82fb120796a,
subnet-06d7f09cca266494f,
subnet-04fef07ed760f4ba5,

Step 2: Choose namespace

Edit

Namespace

Namespace is a collection of database objects and users. Data properties include database name and password, permissions, and encryption and security.

Target namespace
myfirstnamespace

Database name and password

Database name
dev

Admin user credentials
admin

Permissions

Default IAM role
Not applied

IAM roles
[arn:aws:iam::170303796048:role/aws_red_role](#)

Encryption and security

AWS KMS encryption
AWS owned key

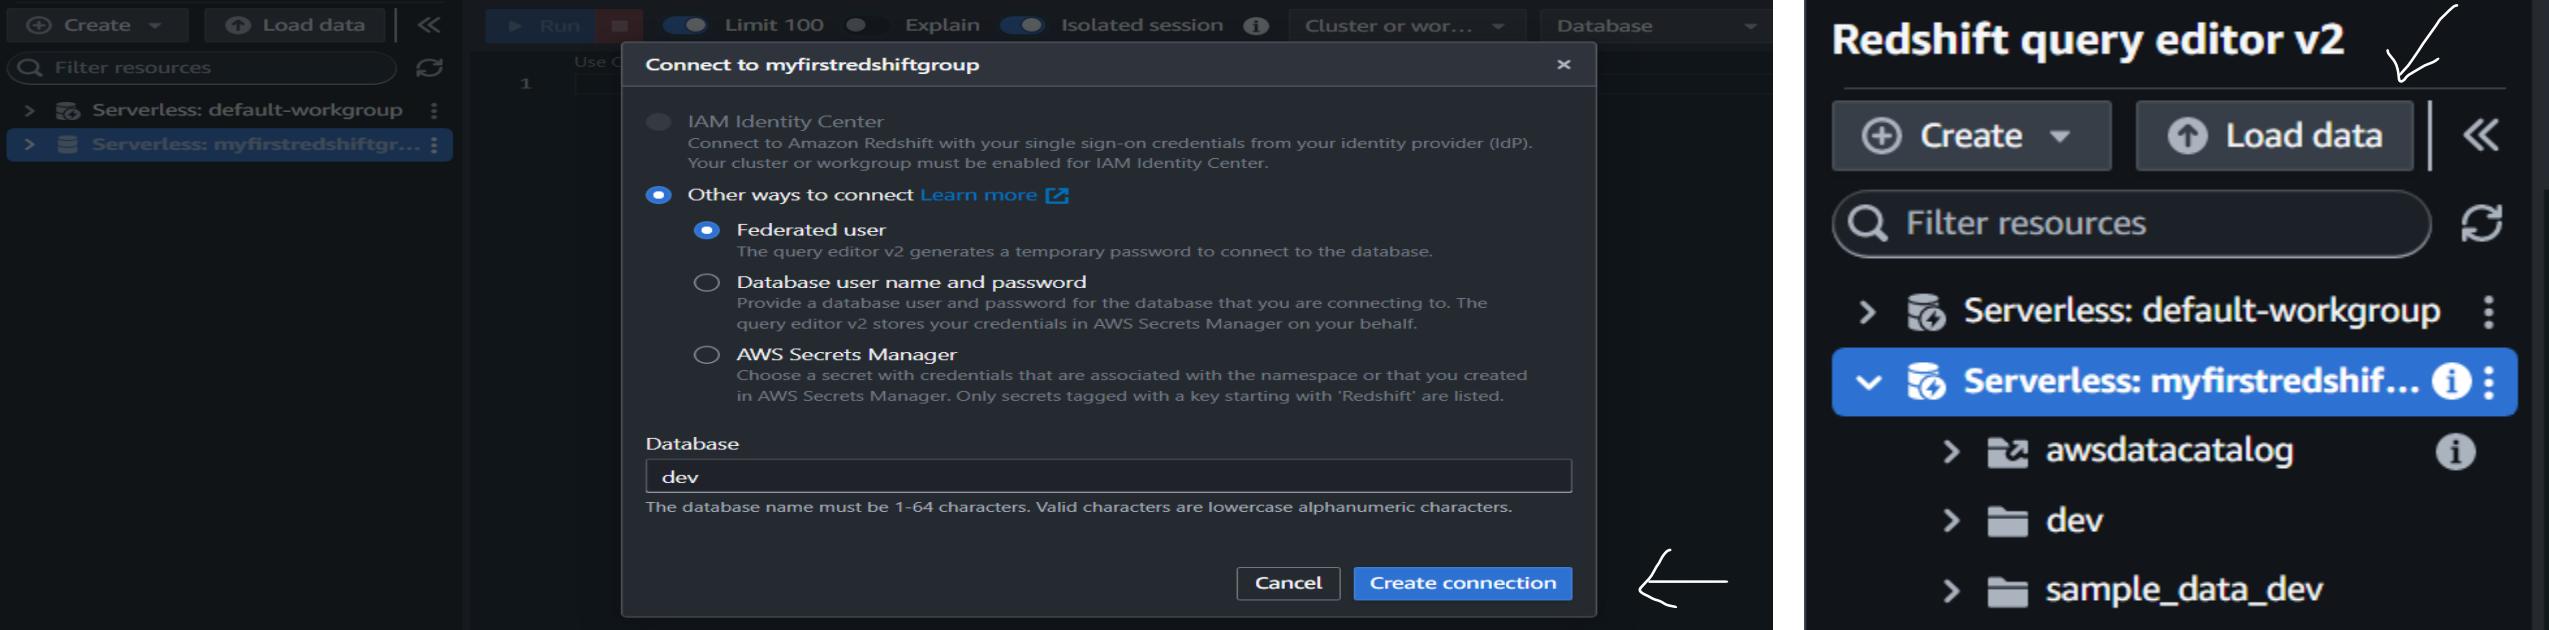
Audit logging
Off

Namespaces / Workgroups [Info](#)

Namespace	Status	Workgroup	Status
default-namespace	Available	default-workgroup	Available
myfirstnamespace	Available	myfirstredshiftgroup	Available

You can go to namespace , workgroup or go on above choose Query data then this interface appeared.

The screenshot shows the AWS Redshift Query Editor v2 interface. The sidebar on the left has icons for Services, Editor (selected), Queries, Notebooks, Charts, History, Scheduled queries, and other AWS services. The main area is titled "Redshift query editor v2". It shows an "Untitled 1" tab with a single line of code "1". The toolbar includes "Run", "Limit 100", "Explain", "Isolated session", "Cluster or wor...", "Database", "Schedule", and other options. A large hand-drawn arrow points upwards from the bottom of the page towards the "Editor" icon in the sidebar.



Load data

Data source

- Load from S3 bucket (selected)
- Load from local file

S3 URI: s3://bucket/prefix/object

Browse S3

S3 file location

Manifest file

File format: CSV

File options >

No compression

Delimiter character: ,

Specifies the single ASCII character that is used to separate fields in the input file, such as a pipe character ('|'), a comma (,), or a tab (\t).

Ignore header rows: 1

Treats the specified number_rows as a file header and doesn't load them. Use this option to skip file headers in all files in a parallel load.

Advanced settings

Data conversion parameters >

Load operations >

Cancel Next

Choose archive in S3

S3 buckets

Name	Creation date
elasticbeanstalk-eu-north-1-170303796048	November 29, 2023, 14:36:39 (UTC+05:30)
redshi-bucket	November 28, 2023, 15:03:54 (UTC+05:30)

Choose archive in S3

S3 buckets > redshi-bucket

Objects

Key	Last modified	Size
currency.csv	November 29, 2023, 16:14:34 (UTC...)	3.8 KB

1 to 1 of 1 Page 1 of 1

Choose

Load data

Table options

Load existing table
Load data into an existing table

Load new table
Create table with detected schema

Cluster or workgroup: Serverless: myfirstredshiftgroup Database: dev Schema: public Table: currency

IAM role: arn:aws:iam::170303796048:role/aws_red_role

Columns **Table details**

Selected S3 file: currency.csv

+ Add column

Column name	Data type	Encoding
col0	VARCHAR	No selection
col1	VARCHAR	No selection
col2	VARCHAR	No selection

Column options

Default value:

- Custom
- Empty string
- NULL
- No default value

Size:

Create table

Load data

Data source

Load from S3 bucket

S3 URI: s3://redshi-bucket/currency.csv

eu-north-1

Manifest file

File format: CSV

File options: >

No compression: >

Delimiter character: ,

Specifies the single ASCII character that is used to separate fields in the input file, such as a pipe character ('|'), a comma (,), or a tab (\t).

Ignore header rows: 1

Treats the specified number_rows as a file header and doesn't load them. Use this option to skip file headers in all files in a parallel load.

Advanced settings

Data conversion parameters: >

Load operations: >

Load data

Review an currency table is created successfully.

Cluster or workgroup: Serverless: myfirstredshiftgroup Database: dev Schema: public Table: currency

IAM role: arn:aws:iam::170303796048:role/aws_red_role

Table summary

Column name	Data type	Encoding
col0	VARCHAR	No selection
col1	VARCHAR	No selection
col2	VARCHAR	No selection

Redshift query editor v2

+ Untitled 1 currency table is created successfully.

Run Limit 100 Explain Isolated session Serverless: m... dev Schedule Load data Create Filter resources

COPY dev.public.currency FROM 's3://redshi-bucket/currency.csv' IAM_ROLE 'arn:aws:iam::170303796048:role/aws_red_role' FORMAT AS CSV
DELIMITER ',' QUOTE '\"' IGNOREHEADER 1 REGION AS 'eu-north-1'

Serverless: default-workgroup Serverless: myfirstredshif... i:

- awsdatacatalog i
- dev
- sample_data_dev

Result 1 Export Chart

Summary

Info:

- Load into table 'currency' completed, 163 record(s) loaded successfully.

Returned rows: 0
Elapsed time: 1.8s
Result set query:

```
/* RQEY2-ACEImhYATy */
COPY dev.public.currency FROM 's3://redshi-bucket/currency.csv' IAM_ROLE 'arn:aws:iam::170303796048:role/aws_red_role' FORMAT AS CSV DELIMITER ',' QUOTE '\"' IGNOREHEADER 1
```

hell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Under dev- public – tables – you will find uploaded file.

Redshift query editor v2

Editor

Queries

Notebooks

Charts

History

Scheduled queries

Untitled 1

currency table is created successfully.

Run Explain Isolated session Serverless: m... dev Schedule

COPY dev.public.currency FROM 's3://redshi-bucket/currency.csv' IAM_ROLE 'arn:aws:iam::170303796048:role/aws_red_role' FORMAT AS CSV
DELIMITER ',' QUOTE '\"' IGNOREHEADER 1 REGION AS 'eu-north-1'

Serverless: default-workgroup

Serverless: myfirstredshift... i:

awsdatacatalog i

dev

pg_auto_copy

public

Tables 1 currency ✓

Views 0

Functions 0

Stored proc... 0

sample_data_dev

Result 1

Export Chart

Summary

Info:

- Load into table 'currency' completed, 163 record(s) loaded successfully.

Returned rows: 0
Elapsed time: 1.8s
Result set query:

```
/* RQEY2-ACEImhYATy */  
COPY dev.public.currency FROM 's3://redshi-bucket/currency.csv' IAM_ROLE 'arn:aws:iam::170303796048:role/aws_red_role' FORMAT AS CSV DELIMITER ',' QUOTE '\"' IGNOREHEADER 1
```

The screenshot shows the Redshift Query Editor interface. On the left is a sidebar with navigation links: Editor, Queries, Notebooks, Charts, History, and Scheduled queries. The main area has tabs for Untitled 1 and currency table is created successfully. The currency table is highlighted. Below the tabs are buttons for Run, Explain, Isolated session, and Serverless. The current database is set to dev. The main pane displays a single query: a COPY command to create a table named currency from an S3 bucket. The table was created successfully, as indicated by the green success message at the top. The summary section shows 163 records loaded and an elapsed time of 1.8 seconds. The result set is empty, and the full query is shown at the bottom.

Double click the currency file and click RUN.

The screenshot shows the Redshift query editor v2 interface. On the left, there's a sidebar with icons for Editor, Queries, Notebooks, Charts, History, Scheduled queries, and a gear icon. The main area has tabs for 'Untitled' and 'Redshift query editor v2'. In the Untitled tab, a green success message says 'currency table is created successfully.' Below it, there are buttons for Run, Limit 100, Explain, Isolated session, and Serverless: m...'. A dropdown shows 'dev'. To the right are buttons for Schedule, Save, and more. The Redshift query editor v2 tab shows a query: 'SELECT * FROM "dev"."public"."currency";'. The results pane shows a table named 'currency' with 100 rows. The columns are 'Field', 'Type', 'NL', 'CMP', 'col0', 'col1', and 'col2'. The data includes entries like United Arab Emirates dirham, Afghan afghani, Albanian lek, Armenian dram, Netherlands Antillean guilder, Angolan kwanza, Argentine peso, Australian dollar, and Aruban florin. At the bottom, it says 'Elapsed time: 143 ms Total rows: 100'. The footer includes CloudShell, Feedback, Copyright 2023, Privacy, Terms, and Cookie preferences.

Field	Type	NL	CMP	col0	col1	col2
A	col0	character varying(256)	NULL	Izo	AED	United Arab Emirates dirham
A	col1	character varying(256)	NULL	Izo	AFN	Afghan afghani
A	col2	character varying(256)	NULL	Izo	ALL	Albanian lek
					AMD	Armenian dram
					ANG	Netherlands Antillean guilder
					AOA	Angolan kwanza
					ARS	Argentine peso
					AUD	Australian dollar
					AWG	Aruban florin

CloudShell Feedback © 2023, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

If we want to run certain column then insert column name at first and click RUN.

The screenshot shows a database interface with a dark theme. At the top, there is a message bar with a green background and white text: "currency table is created successfully." Below the message bar, the main area contains a query editor window. The query in the editor is:

```
1  SELECT Col1 FROM "dev"."public"."currency";
```

Below the editor, a results table is displayed with the title "Result 1 (100)". The table has one column labeled "col1". The data rows are:

col1
J\$
đ
L
AMD
f
Kz
\$
\$
₹

At the bottom right of the results table, there are buttons for "Export" and "Chart". The status bar at the bottom of the interface shows "Elapsed time: 103 ms" and "Total rows: 100".